

# AI System & Agent Design Principles with API Error Handling Architecture

This document is a concise, language-agnostic reference for building secure, stable, and maintainable AI-driven systems. It extends the core design principles with a standardized API Error Handling Architecture to ensure correctness, auditability, and long-term reliability.

## 1. Secure by Design

Systems and agents must be built with safety and predictability as defaults.

- Least Privilege: Grant only the minimum required capabilities.
- Explicit Boundaries: Every tool has defined inputs, outputs, and side effects.
- No Implicit Trust: Validate all incoming and outgoing data.
- Auditability: Log tool usage, errors, and state changes.
- Idempotent Actions: Retries must not duplicate or corrupt data.
- Fixed Capability Set: Agents cannot create or extend their own capabilities.

## 2. SOLID Principles Applied to AI Systems

- Single Responsibility: Each tool or module does one thing.
- Open/Closed: Extend via new components, not modifications.
- Liskov Substitution: Swap implementations without breaking consumers.
- Interface Segregation: Small, task-specific interfaces only.
- Dependency Inversion: Depend on stable contracts, not implementations.

## 3. DRY (Don't Repeat Yourself)

Avoid duplicating validation, schemas, error handling, or business logic. Centralization increases reliability, security, and maintainability.

## 4. YAGNI (You Aren't Gonna Need It)

Do not build speculative capabilities. Add tools only when a real use case exists and keep scopes small and intentional.

## 5. KISS / Law of Simplicity

Simple systems fail in predictable ways. Keep data flow explicit, reduce hidden branching, and avoid nested orchestration logic.

## 6. Replaceability & Modularity

LLMs, vector stores, memory layers, tools, and planning modules must be swappable through clear contracts without system-wide refactors.

## 7. Contracts Everywhere

Use structured, enforced schemas at every boundary to reduce improvisation, standardize behavior, and increase predictability and safety.

## 8. API Error Handling Architecture (Language-Agnostic)

Error handling is treated as a first-class architectural concern, not ad-hoc logic. Errors are part of the domain and enforced consistently across all system boundaries.

- Domain Truth: Define canonical error codes, categories, and semantics independent of transport, frameworks, or databases.
- Boundary Translation: Map domain errors and unknown failures into transport-specific responses (HTTP, jobs, webhooks) without inventing new meaning.
- Enforced Consistency: All entry points flow through a single choke point. No custom error shapes, no duplicated try/catch logic.

This architecture ensures predictable failures, auditability, and makes AI-generated code operate inside strict guardrails instead of inventing behavior.

## 9. Practical Checklist

- Does this component do exactly one thing?
- Is there a clear input/output contract?
- Is this capability necessary right now (YAGNI)?
- Does it follow least privilege?
- Is logic duplicated anywhere?
- Can the implementation be swapped safely?
- Does failure remain predictable and auditable?

## Core Mindset

Design AI systems as adversarial until proven otherwise. Keep components small, contracts explicit, errors centralized, and behavior predictable. This is how you build AI systems that behave.